

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=pd.read_csv("AirQuality2.csv",sep=';')
```

```
df.head()
```

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)
PT08.S2(NMHC)	\					
0	10/03/2004	18.00.00	2,6	1360.0	150.0	11,9
1046.0						
1	10/03/2004	19.00.00	2	1292.0	112.0	9,4
955.0						
2	10/03/2004	20.00.00	2,2	1402.0	88.0	9,0
939.0						
3	10/03/2004	21.00.00	2,2	1376.0	80.0	9,2
948.0						
4	10/03/2004	22.00.00	1,6	1272.0	51.0	6,5
836.0						

	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)	T
RH	\					
0	166.0	1056.0	113.0	1692.0	1268.0	13,6
48,9						
1	103.0	1174.0	92.0	1559.0	972.0	13,3
47,7						
2	131.0	1140.0	114.0	1555.0	1074.0	11,9
54,0						
3	172.0	1092.0	122.0	1584.0	1203.0	11,0
60,0						
4	131.0	1205.0	116.0	1490.0	1110.0	11,2
59,6						

	AH	Unnamed: 15	Unnamed: 16
0	0,7578	NaN	NaN
1	0,7255	NaN	NaN
2	0,7502	NaN	NaN
3	0,7867	NaN	NaN
4	0,7888	NaN	NaN

```
df.shape
```

```
(9471, 17)
```

```
#data cleaning
```

```
df=df.drop_duplicates()
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 9358 entries, 0 to 9357
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  9357 non-null   object
1   Time                  9357 non-null   object
2   CO(GT)                9357 non-null   object
3   PT08.S1(CO)           9357 non-null   float64
4   NMHC(GT)              9357 non-null   float64
5   C6H6(GT)              9357 non-null   object
6   PT08.S2(NMHC)         9357 non-null   float64
7   NOx(GT)               9357 non-null   float64
8   PT08.S3(NOx)          9357 non-null   float64
9   NO2(GT)               9357 non-null   float64
10  PT08.S4(NO2)          9357 non-null   float64
11  PT08.S5(O3)           9357 non-null   float64
12  T                      9357 non-null   object
13  RH                     9357 non-null   object
14  AH                     9357 non-null   object
15  Unnamed: 15            0 non-null      float64
16  Unnamed: 16            0 non-null      float64
dtypes: float64(10), object(7)
memory usage: 1.3+ MB
```

```
df.isna().sum()
```

```
Date                1
Time                1
CO(GT)              1
PT08.S1(CO)         1
NMHC(GT)            1
C6H6(GT)            1
PT08.S2(NMHC)       1
NOx(GT)             1
PT08.S3(NOx)        1
NO2(GT)             1
PT08.S4(NO2)        1
PT08.S5(O3)         1
T                   1
RH                  1
AH                  1
Unnamed: 15         9358
Unnamed: 16         9358
dtype: int64
```

```
df=df.drop(['Unnamed: 15','Unnamed: 16'],axis=1)
```

```
df
```

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	\
0	10/03/2004	18.00.00	2,6	1360.0	150.0	11,9	
1	10/03/2004	19.00.00	2	1292.0	112.0	9,4	
2	10/03/2004	20.00.00	2,2	1402.0	88.0	9,0	
3	10/03/2004	21.00.00	2,2	1376.0	80.0	9,2	
4	10/03/2004	22.00.00	1,6	1272.0	51.0	6,5	
...	...	...	...	...	...	...	
9353	04/04/2005	11.00.00	2,4	1163.0	-200.0	11,4	
9354	04/04/2005	12.00.00	2,4	1142.0	-200.0	12,4	
9355	04/04/2005	13.00.00	2,1	1003.0	-200.0	9,5	
9356	04/04/2005	14.00.00	2,2	1071.0	-200.0	11,9	
9357	NaN	NaN	NaN	NaN	NaN	NaN	

	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	\
0	1046.0	166.0	1056.0	113.0	1692.0	
1	955.0	103.0	1174.0	92.0	1559.0	
2	939.0	131.0	1140.0	114.0	1555.0	
3	948.0	172.0	1092.0	122.0	1584.0	
4	836.0	131.0	1205.0	116.0	1490.0	
...	...	...	...	...	...	
9353	1027.0	353.0	604.0	179.0	1264.0	
9354	1063.0	293.0	603.0	175.0	1241.0	
9355	961.0	235.0	702.0	156.0	1041.0	
9356	1047.0	265.0	654.0	168.0	1129.0	
9357	NaN	NaN	NaN	NaN	NaN	

	PT08.S5(O3)	T	RH	AH
0	1268.0	13,6	48,9	0,7578
1	972.0	13,3	47,7	0,7255
2	1074.0	11,9	54,0	0,7502
3	1203.0	11,0	60,0	0,7867
4	1110.0	11,2	59,6	0,7888
...	...	...	...	...
9353	1269.0	24,3	23,7	0,7119
9354	1092.0	26,9	18,3	0,6406
9355	770.0	28,3	13,5	0,5139
9356	816.0	28,5	13,1	0,5028
9357	NaN	NaN	NaN	NaN

[9358 rows x 15 columns]

```

df=df.rename(columns={'T': 'Temperature'})
df=df.rename(columns={'AH': 'Absolute Humidity'})
df=df.rename(columns={'RH': 'Relative Humidity'})
df

```


	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	\
0	10/03/2004	18.00.00	2,6	1360.0	150.0	11,9	
1	10/03/2004	19.00.00	2	1292.0	112.0	9,4	
2	10/03/2004	20.00.00	2,2	1402.0	88.0	9,0	

3	10/03/2004	21.00.00	2,2	1376.0	80.0	9,2
4	10/03/2004	22.00.00	1,6	1272.0	51.0	6,5
...	...	...	...	...	...	...
9353	04/04/2005	11.00.00	2,4	1163.0	-200.0	11,4
9354	04/04/2005	12.00.00	2,4	1142.0	-200.0	12,4
9355	04/04/2005	13.00.00	2,1	1003.0	-200.0	9,5
9356	04/04/2005	14.00.00	2,2	1071.0	-200.0	11,9
9357	NaN	NaN	NaN	NaN	NaN	NaN

	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	\
0	1046.0	166.0	1056.0	113.0	1692.0	
1	955.0	103.0	1174.0	92.0	1559.0	
2	939.0	131.0	1140.0	114.0	1555.0	
3	948.0	172.0	1092.0	122.0	1584.0	
4	836.0	131.0	1205.0	116.0	1490.0	
...	...	...	...	...	...	...
9353	1027.0	353.0	604.0	179.0	1264.0	
9354	1063.0	293.0	603.0	175.0	1241.0	
9355	961.0	235.0	702.0	156.0	1041.0	
9356	1047.0	265.0	654.0	168.0	1129.0	
9357	NaN	NaN	NaN	NaN	NaN	

	PT08.S5(03)	Temperature	Relative Humidity	Absolute Humidity
0	1268.0	13,6	48,9	0,7578
1	972.0	13,3	47,7	0,7255
2	1074.0	11,9	54,0	0,7502
3	1203.0	11,0	60,0	0,7867
4	1110.0	11,2	59,6	0,7888
...	...	...	...	...
9353	1269.0	24,3	23,7	0,7119
9354	1092.0	26,9	18,3	0,6406
9355	770.0	28,3	13,5	0,5139
9356	816.0	28,5	13,1	0,5028
9357	NaN	NaN	NaN	NaN

[9358 rows x 15 columns]

```
df=df.replace(',','.',regex=True)
```

```
float_col=["CO(GT)","C6H6(GT)","Temperature","Relative Humidity","Absolute Humidity"]
```

```
df[float_col]=df[float_col].astype(float)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 9358 entries, 0 to 9357
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----

0	Date	9357	non-null	object
1	Time	9357	non-null	object
2	CO(GT)	9357	non-null	float64
3	PT08.S1(CO)	9357	non-null	float64
4	NMHC(GT)	9357	non-null	float64
5	C6H6(GT)	9357	non-null	float64
6	PT08.S2(NMHC)	9357	non-null	float64
7	NOx(GT)	9357	non-null	float64
8	PT08.S3(NOx)	9357	non-null	float64
9	NO2(GT)	9357	non-null	float64
10	PT08.S4(NO2)	9357	non-null	float64
11	PT08.S5(O3)	9357	non-null	float64
12	Temperature	9357	non-null	float64
13	Relative Humidity	9357	non-null	float64
14	Absolute Humidity	9357	non-null	float64

dtypes: float64(13), object(2)
memory usage: 1.1+ MB

```
df=df.drop(['Date','Time'],axis=1)
```

```
df.isna().sum()
```

CO(GT)	1
PT08.S1(CO)	1
NMHC(GT)	1
C6H6(GT)	1
PT08.S2(NMHC)	1
NOx(GT)	1
PT08.S3(NOx)	1
NO2(GT)	1
PT08.S4(NO2)	1
PT08.S5(O3)	1
Temperature	1
Relative Humidity	1
Absolute Humidity	1

dtype: int64

```
df=df.dropna()
```

```
df.isna().sum()
```

CO(GT)	0
PT08.S1(CO)	0
NMHC(GT)	0
C6H6(GT)	0
PT08.S2(NMHC)	0
NOx(GT)	0
PT08.S3(NOx)	0
NO2(GT)	0
PT08.S4(NO2)	0
PT08.S5(O3)	0

```

Temperature      0
Relative Humidity 0
Absolute Humidity 0
dtype: int64

```

```
df.head()
```

	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	\
0	2.6	1360.0	150.0	11.9	1046.0	166.0	
1	2.0	1292.0	112.0	9.4	955.0	103.0	
2	2.2	1402.0	88.0	9.0	939.0	131.0	
3	2.2	1376.0	80.0	9.2	948.0	172.0	
4	1.6	1272.0	51.0	6.5	836.0	131.0	

	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)	Temperature	\
0	1056.0	113.0	1692.0	1268.0	13.6	
1	1174.0	92.0	1559.0	972.0	13.3	
2	1140.0	114.0	1555.0	1074.0	11.9	
3	1092.0	122.0	1584.0	1203.0	11.0	
4	1205.0	116.0	1490.0	1110.0	11.2	

	Relative Humidity	Absolute Humidity
0	48.9	0.7578
1	47.7	0.7255
2	54.0	0.7502
3	60.0	0.7867
4	59.6	0.7888

```
#data integration
```

```
common_col=['CO(GT)', 'NO2(GT)']
```

```
data1=df[common_col+['C6H6(GT)', 'NOx(GT)']]
```

```
data2=df[common_col+
```

```
["PT08.S1(CO)", "PT08.S2(NMHC)", "PT08.S3(NOx)", "PT08.S4(NO2)", "PT08.S5(O3)", "Temperature", "Relative Humidity", "Absolute Humidity"]]
```

```
data1_s=data1.head(100)
```

```
data2_s=data2.head(100)
```

```
inner_merged=pd.merge(data1_s, data2_s, on=common_col, how="inner")
```

```
inner_merged.head()
```

	CO(GT)	NO2(GT)	C6H6(GT)	NOx(GT)	PT08.S1(CO)	PT08.S2(NMHC)	\
0	2.6	113.0	11.9	166.0	1360.0	1046.0	
1	2.0	92.0	9.4	103.0	1292.0	955.0	
2	2.2	114.0	9.0	131.0	1402.0	939.0	
3	2.2	122.0	9.2	172.0	1376.0	948.0	
4	1.6	116.0	6.5	131.0	1272.0	836.0	

	PT08.S3(NOx)	PT08.S4(NO2)	PT08.S5(O3)	Temperature	Relative Humidity	\
--	--------------	--------------	-------------	-------------	-------------------	---

0	1056.0	1692.0	1268.0	13.6
48.9				
1	1174.0	1559.0	972.0	13.3
47.7				
2	1140.0	1555.0	1074.0	11.9
54.0				
3	1092.0	1584.0	1203.0	11.0
60.0				
4	1205.0	1490.0	1110.0	11.2
59.6				

Absolute Humidity	
0	0.7578
1	0.7255
2	0.7502
3	0.7867
4	0.7888

```
left_merged=pd.merge(data1_s,data2_s,on=common_col,how='left')
left_merged.head()
```


	CO(GT)	N02(GT)	C6H6(GT)	N0x(GT)	PT08.S1(CO)	PT08.S2(NMHC)	\
0	2.6	113.0	11.9	166.0	1360.0	1046.0	
1	2.0	92.0	9.4	103.0	1292.0	955.0	
2	2.2	114.0	9.0	131.0	1402.0	939.0	
3	2.2	122.0	9.2	172.0	1376.0	948.0	
4	1.6	116.0	6.5	131.0	1272.0	836.0	

	PT08.S3(N0x)	PT08.S4(N02)	PT08.S5(O3)	Temperature	Relative Humidity	\
0	1056.0	1692.0	1268.0	13.6		
48.9						
1	1174.0	1559.0	972.0	13.3		
47.7						
2	1140.0	1555.0	1074.0	11.9		
54.0						
3	1092.0	1584.0	1203.0	11.0		
60.0						
4	1205.0	1490.0	1110.0	11.2		
59.6						

Absolute Humidity	
0	0.7578
1	0.7255
2	0.7502
3	0.7867
4	0.7888

```
#error correction
sns.boxplot(df)
plt.xticks(rotation=45)

([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12],
 [Text(0, 0, 'CO(GT)'),
  Text(1, 0, 'PT08.S1(CO)'),
  Text(2, 0, 'NMHC(GT)'),
  Text(3, 0, 'C6H6(GT)'),
  Text(4, 0, 'PT08.S2(NMHC)'),
  Text(5, 0, 'NOx(GT)'),
  Text(6, 0, 'PT08.S3(NOx)'),
  Text(7, 0, 'NO2(GT)'),
  Text(8, 0, 'PT08.S4(NO2)'),
  Text(9, 0, 'PT08.S5(O3)'),
  Text(10, 0, 'Temperature'),
  Text(11, 0, 'Relative Humidity'),
  Text(12, 0, 'Absolute Humidity')])
```




```

def remove_outliers(col):
    Q1=col.quantile(0.25)
    Q3=col.quantile(0.75)
    IQR=Q3-Q1
    threshold=2.5*IQR
    outlier_mask=(col<Q1-threshold) | (col>Q3+threshold)
    return col[~outlier_mask]

df.columns


Index(['C0(GT)', 'PT08.S1(CO)', 'NMHC(GT)', 'C6H6(GT)',
      'PT08.S2(NMHC)',
      'NOx(GT)', 'PT08.S3(NOx)', 'NO2(GT)', 'PT08.S4(NO2)',
      'PT08.S5(O3)',
      'Temperature', 'Relative Humidity', 'Absolute Humidity'],
      dtype='object')

num_cols=df.select_dtypes(include=['number']).columns

for col in num_cols:
    cleaned_col=remove_outliers(df[col])
    df.loc[cleaned_col.index,col]=cleaned_col

sns.boxplot(df)
plt.xticks(rotation=45)
plt.show()

```



```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler=MinMaxScaler()
```

```
df[num_cols]=scaler.fit_transform(df[num_cols])
```

```
df.head()
```

	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)
0	0.956111	0.696429	0.251980	0.803565	0.516156	0.217987
1	0.953280	0.666071	0.224622	0.794084	0.478459	0.180465
2	0.954224	0.715179	0.207343	0.792567	0.471831	0.197141
3	0.954224	0.703571	0.201584	0.793326	0.475559	0.221560
4	0.951392	0.657143	0.180706	0.783087	0.429163	0.197141

	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)	Temperature
0	0.435657	0.579630	0.635966	0.539111	0.873262

1	0.476587	0.540741	0.591261	0.430408	0.872036
2	0.464794	0.581481	0.589916	0.467866	0.866312
3	0.448144	0.596296	0.599664	0.515241	0.862633
4	0.487340	0.585185	0.568067	0.481087	0.863451

	Relative Humidity	Absolute Humidity
0	0.862141	0.992715
1	0.857984	0.992556
2	0.879806	0.992678
3	0.900589	0.992858
4	0.899203	0.992869

#data model building

y=df['Temperature']

X=df.drop('Temperature',axis=1)

from sklearn.model_selection import train_test_split

from sklearn.metrics import

r2_score,mean_absolute_error,mean_squared_error

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)

from sklearn.linear_model import LinearRegression

model1=LinearRegression()

model1.fit(X_train,y_train)

y_pred=model1.predict(X_test)

print("r2 score:",r2_score(y_test,y_pred))

print("mean absolute error:",mean_absolute_error(y_test,y_pred))

print("mean squared error:",mean_squared_error(y_test,y_pred))

print("accuracy:",r2_score(y_test,y_pred)*100)

r2 score: 0.995077596394832

mean absolute error: 0.009681121939987306

mean squared error: 0.00015160045647323898

accuracy: 99.5077596394832

from sklearn.tree import DecisionTreeRegressor

model2=DecisionTreeRegressor()

model2.fit(X_train,y_train)

y_pred=model2.predict(X_test)

print("r2 score:",r2_score(y_test,y_pred))

print("mean absolute error:",mean_absolute_error(y_test,y_pred))

print("mean squared error:",mean_squared_error(y_test,y_pred))

print("accuracy:",r2_score(y_test,y_pred)*100)

r2 score: 0.9998792916282065

mean absolute error: 0.001236319894332979

mean squared error: 3.717583061418078e-06

accuracy: 99.98792916282065

```
from sklearn.ensemble import RandomForestRegressor
model3 = RandomForestRegressor()

from sklearn.model_selection import cross_val_score
scores = cross_val_score(model3, X, y, cv=5, scoring='r2')
print("Mean R2 from Cross-validation:", scores.mean())
```